

Database Management System (DBMS)

Applications

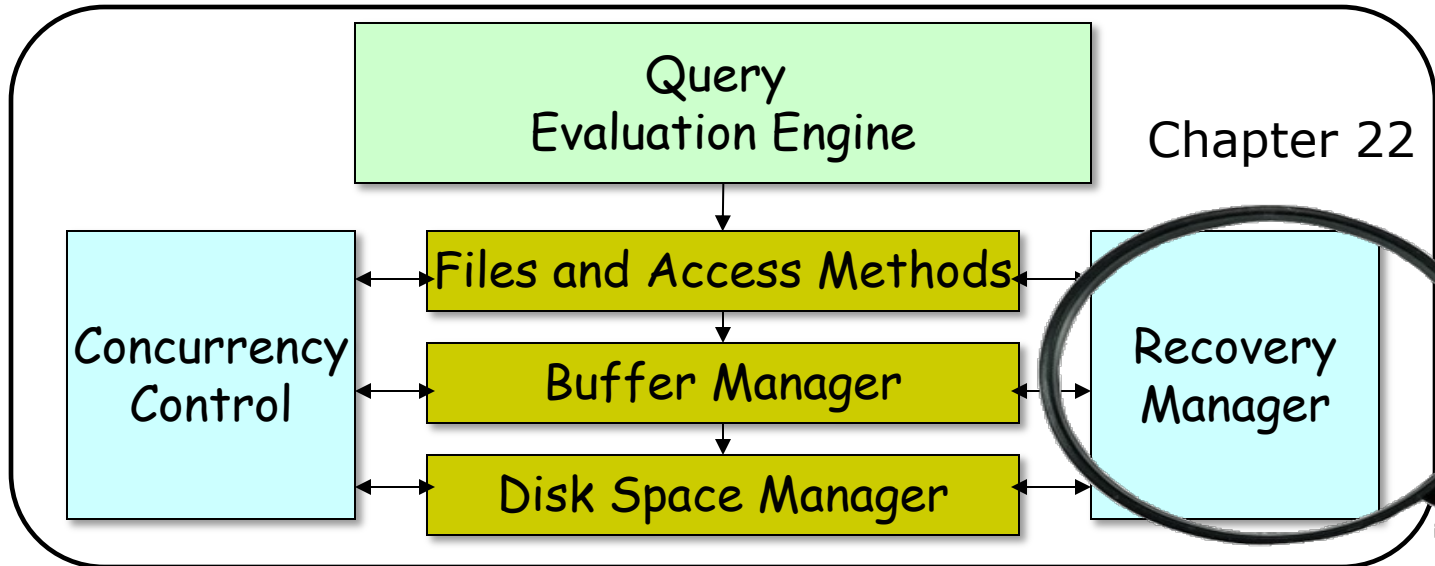
Web
Forms

Embedded
SQL

Interactive
SQL

SQL Commands

DBMS



Database

Data

Indexes

Catalog

Many Things Can Go Wrong

- ☐ **Interference with other concurrent activities**
- ☐ User may decide to interrupt the program
- ☐ Account number does not exist
- ☐ Integer overflow
- ☐ disk head crash, system goes down
- ☐ Buffer congestion
- ☐ Error during data transfer
- ☐ Power failure



ACID Properties

Property	Dealt with by
A, D	Recovery Techniques
I	Concurrency Control Techniques
C	Checks, Assertions, Triggers Applications Programmers

Purpose of Database Recovery

- ❑ To bring the database into the **last consistent state**, which existed prior to the failure
- ❑ To preserve transaction properties:

- **Atomicity & Durability**



- ❑ Example: System crashes before a fund transfer transaction completes its execution
 - The database must be restored to the state before the transaction modified any of the accounts.

Types of Failures

- Transaction failure:

- Incorrect input, deadlock, etc.

- System failure:

- Addressing error, application error, operating system fault, RAM failure, etc.

- Media failure:

- Disk head crash, power disruption, etc.

Motivation


❑ Atomicity:

- Transactions may abort (“Rollback”)

❑ Durability:

- What if DBMS stops running?

- ❖ Desired Behavior after system restarts:
 - T1, T2 & T3 should be durable
 - T4 & T5 should be aborted (effects not seen)



The Goals of Recovery

1. When a transaction *T commits*

- Make the updates **permanent** in the database so that they can survive subsequent failures

2. When a transaction *T aborts*

- Obliterate any updates on data items by **aborted transaction** in the database
- Obliterate the effects of *T* **on other transactions**; i.e., transactions that read data items updated by *T*

3. When the system *crashes* after a system failure

- Bring the database to its most recent consistent state

Recovery Actions

- Recovery protocols implement **two actions**:
 - **Undo** action:
 - Required for **atomicity**
 - Undoes all updates on the stable storage by an uncommitted transaction
 - **Redo** action:
 - Required for **durability**
 - Redoes the update on the stable storage of committed transaction

Logging

- A **log** or *journal* is a sequence of records which represent all modifications to the database in the order in which they actually occurred
- Log records may describe either *physical* changes or *logical* database operations
 - A **physical log** contains information about the actual values of data items written by transactions.
 - State before change, **before image**
 - State after change, **after image**
 - Transition causing the change
 - A **logical log** represents higher level operations; e.g., insert a record

Transaction Log

Types of log record:

1. **[start_transaction,T]**: Records that transaction T has started execution
2. **[write_item,T,X,old_value,new_value]**: Records that transaction T has changed the value of database item X from old_value to new_value.
3. **[read_item,T,X]**: Records that transaction T has read the value of database item X
4. **[commit,T]**: Records that transaction T has completed successfully
5. **[abort,T]**: Records that transaction T has been aborted

Transaction Log

TID	Back P	Next P	Operation	Data Item	BFIM	AFIM
T1	0	2	Begin			
T1	1	4	W	X	X = 100	X = 200
T2	0	7	Begin			
T1	2	5	W	Y	Y = 50	Y = 100
T1	4	8	R	M	M = 200	M = 200
T3	0	9	Begin			
T2	3	10	W	N	N = 20	N = 10
T1	5	nil	End			
T3	6	11	W	N	N = 10	N = 30

- **BFIM**: old values before modification (BeFore IMage)
- **AFIM**: new value after modification (AFter IMage)
- **Back P** and **Next P** point to the previous and next log records of the same transaction

What & How to Recover?

☐ What & how to **undo**?

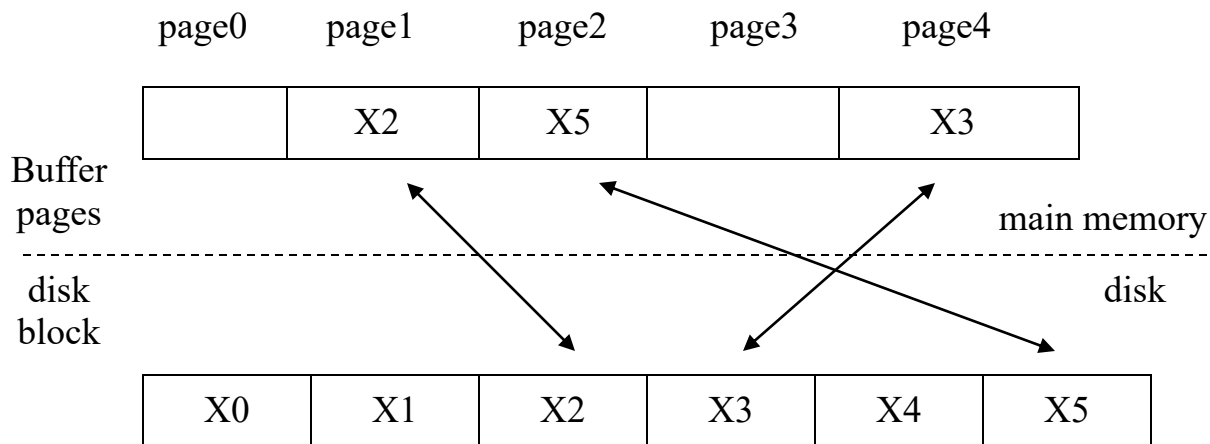
- For all uncommitted transactions
- Scan the log from end to beginning
- Use the BFIM

☐ What & how to **redo**?

- For all committed transactions
- Scan the log from beginning to end
- Use the AFIM

Buffer Manager

- ❑ The goal of Cache or Buffer Manager is to **maximize the likelihood** that a block of data needed by a transaction is in main memory
- ❑ The main memory is **partitioned** into *buffer blocks* or *buffer pages* (or simply buffers, blocks or pages)
- ❑ The size of a page is equal to the *disk block* size



Cache Flushing

- **Steal:** A cache page can be flushed before transaction commits – when?
- **No-Steal:** A cache cannot be flushed before transaction commits
- **Force:** Cache is flushed to disk when transaction commits
- **No-Force:** Otherwise! So when is it flushed?

Cache Flushing

- Leads to four different ways for handling recovery:
 - Steal/No-Force (Undo/Redo),
 - Steal/Force (Undo/No-redo),
 - No-Steal/No-Force (No-undo/Redo) and
 - No-Steal/Force (No-undo/No-redo).

Simplest Recovery Manager

□ Force

- No redo of committed transactions

□ No-Steal

- No undo of incomplete transactions

	No Steal	Steal
Force	Simple	
No Force		

Simplest Recovery Manager - Problems

☐ Force

- High I/O costs!

☐ No-Steal

- Large cache!

☐ Practical Recovery

☐ No-Force

- Need to Redo

☐ Steal

- Need to Undo

	No Steal	Steal
Force	Simple	
No Force		Practical

Write-Ahead Logging (WAL)

- Log is necessary for recovery and it must be available to recovery manager
- The **Write-Ahead Logging** (WAL) protocol states that:
- **For Undo:**
 - Before a data item's AFIM is flushed to the database disk (overwriting the BFIM) its **BFIM must be written** to the log and the log must be saved on a stable store (log disk)
- **For Redo:**
 - Before a transaction executes its **commit** operation, all its **AFIMs must be written** to the log and the log must be saved on a stable store

Commit Point of a Transaction

- ❑ A transaction T reaches its **commit point** when:
 1. All its operations have been executed successfully
 2. The effect of all the transaction operations on the database has been **recorded in the log**.
- ❑ Beyond the commit point, its effect is assumed to be *permanently recorded* in the database.
- ❑ The transaction then writes an entry [commit,T] into the log.
- ❑ **Rollback of transactions:** Needed for transactions that have a [start_transaction,T] entry into the log but **no commit entry** [commit,T]

Recall: What & How to Recover?

□ What & how to **undo**?

- For all uncommitted transactions
- Scan the log from end to beginning
- Use the BFIM

□ What & how to **redo**?

- For all committed transactions
- Scan the log from beginning to end
- Use the AFIM

Checkpoint

❑ Scanning the whole log is time-consuming!

❑ Checkpoint

■ Periodically written to the log

■ All cache pages are flushed to disk

- ❖ What transactions should be undo?
- ❖ What transactions should be redo?

