



THE UNIVERSITY
OF QUEENSLAND
AUSTRALIA

Venue _____

Seat Number _____

Student Number

--	--	--	--	--	--	--	--	--	--

Family Name _____

First Name _____

Semester Two Final Examinations, 2017

This paper is for St Lucia Campus students.

Reading Time: 10 minutes

For Examiner Use Only

Question	Mark
----------	------

During reading time - write only on the rough paper provided

This examination paper will be released to the Library

(No electronic aids are permitted e.g. laptops, phones)

Calculators - Casio FX82 series or UQ approved (labelled)

Materials To Be Supplied To Students:

Instructions To Students:

Additional exam materials (eg. answer booklets, rough paper) will be provided upon request.

Please answer all questions on the examination paper.

For Multiple Choice Questions, please circle a single answer.

Total Marks: 100 (to be scaled down to 60)

[illegible]

Question 1 [4 marks] Which of the following is a false statement about B+ trees?

- A. B+-trees are balanced
- B. non-leaf nodes include direct pointers to data records
- C. insertion of a key can lead to node splitting
- D. deletion of a key can lead to node coalescing

Question 2 [4 marks] Which of the following factors determines the size of a bitmap index on an attribute “X” in relation “R”?

- A. The number of distinct values in “X”
- B. The number of tuples in “R”
- C. The data type for attribute “X”
- D. Answers A & B above
- E. Answers A & B & C

Question 3 [4 marks] Which of the following is a false statement about cost-based query optimization?

- A. It selects a query plan in a shorter time than heuristic-based optimization
- B. It requires estimating the execution cost of a query plan
- C. It selects the query plan with the minimum execution cost
- D. All of the above

Questions 4-5 Consider the following relation:

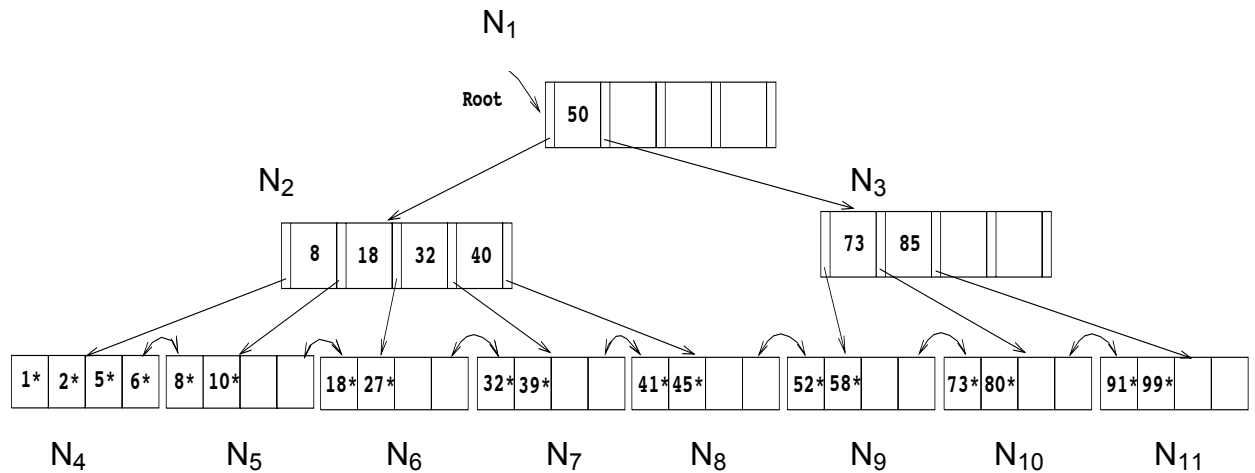
Make	Model	Color	Price
Honda	Accord	Blue	Medium
Honda	Civic	Red	Low
Toyota	Corolla	Black	Low
Toyota	Camry	Red	Medium

Question 4 [4 marks] Assume the relation shown above and a bitmap index is created on attribute '*Price*'. The total number of bits required for that index is:

- A. 2
- B. 4
- C. 8
- D. 12

Question 5 [4 marks] Again, assume the relation shown above, what is the bitmap corresponding to the value '*Honda*'.

- A. 1 1 0 0
- B. 1 1
- C. 1 0 1 0
- D. 0 0



Questions 6-8: Consider a B+ tree index as shown in figure, where index nodes are labeled: N_1, N_2, \dots, N_{11} . Also, assume the following rule applies for redistributing keys after a leaf node split: **Three keys** stay in the old leaf node and the remaining keys move to a new leaf node.

Question 6 [4 marks] Which nodes in the B+ tree index that must be fetched to answer the query: "Get all records with key greater than 30 and less than 75"

- A. $N_1 N_2 N_6 N_7 N_8 N_9 N_{10}$
- B. $N_1 N_2 N_7 N_8 N_9$
- C. $N_1 N_3 N_7 N_8 N_9$
- D. $N_1 N_3 N_7 N_8 N_9 N_{10}$

Question 7 [4 marks] What is the number of leaf nodes after inserting an entry with key "3"?

- A. 8
- B. 9
- C. 11
- D. 12

Question 8 [4 marks] What is the number of non-leaf nodes after inserting an entry with key "3"?

- A. 3
- B. 4
- C. 11
- D. 12

Questions 9-11 Given two relations R1 and R2, where R1 contains N1 tuples, R2 contains N2 tuples, and $N2 > N1 > 0$, answer the following questions.

Question 9 [4 marks] The minimum and maximum number of tuples produced from $R1 \cup R2$ is:

- A. minimum 0, and maximum $N1+N2$
- B. minimum N1, and maximum N2
- C. minimum N1, and maximum $N1+N2$
- D. minimum N2, and maximum $N1+N2$

Question 10 [4 marks] The minimum and maximum number of tuples produced from $R1 \times R2$ is:

- A. minimum 0, and maximum $N1*N2$
- B. minimum N1, and maximum $N1+N2$
- C. minimum $N1*N2$, and maximum $N1*N2$
- D. minimum N2, and maximum $N1+N2$

Question 11 [4 marks] Assume relation R1 contains an attribute named x, the minimum and maximum number of tuples produced from $\sigma_{x=5}(R1)$ is:

- A. minimum 0, and maximum N1
- B. minimum N1, and maximum N1
- C. minimum 1, and maximum N1
- D. minimum N1, and maximum N2

Questions 12-13: Suppose we have two unary (one attribute only) relations, R and S as shown below. Use R for the outer loop and S for the inner loop.

R	S
7	8
2	4
9	2
8	1
3	3
9	2
1	7
3	3
6	

Question 12 [4 marks] Assume a natural join between R and S using **Nested Loop** join (one tuple at a time). The first five results of that join in the **order** that they would be produced by the nested loop is:

- A. 7, 2, 8, 3, 1
- B. 7, 2, 2, 8, 3
- C. 7, 2, 8, 3, 3
- D. None of the above

Question 13 [4 marks] Again, assume a natural join between R and S using **Nested Loop** join (one tuple at a time). The number of iterations needed to finish this join operation is:

- A. 1
- B. 8
- C. 9
- D. 5

Questions 14-15 Consider the relation `Student(Id, Major, Status)`, which has:

- A B+ tree index on `Major` and no other indexes.
- 10,000 tuples of data spread over 100 different blocks.
- The domain of `Status` has 9 values and that of `Major` has 10 different values.

Question 14 [4 marks] What is the estimated number of results returned by the expression $\sigma_{\text{Major}='IT'}(\text{Student})$:

- A. 10
- B. 100
- C. 1,000
- D. 10,000
- E. None of the above

Question 15 [4 marks] What is the selectivity of the expression $\sigma_{\text{Major}='IT' \text{ OR } \text{Major}='CS'}(\text{Student})$:

- A. 0.0
- B. 0.1
- C. 0.2
- D. 1
- E. None of the above

Question 16 [4 marks] Which of the following is a correct statement about transactions?

- A. Redo is needed for atomicity
- B. Undo is needed for durability
- C. Concurrency Control is realized using Triggers and Assertions
- D. Deadlocks do not occur in serial executions
- E. None of the above.

Question 17 [4 marks] Which of the following transaction schedules does not contain conflicting operations? Recall that r = read and w = write.

- A. r1 (A), r2 (A), w1 (C), r1 (B), r2 (B)
- B. r1 (A), r1 (B), w1 (A), r2 (B), r2 (A), w2 (A)
- C. r1 (A), w1 (A), r1 (B), w1 (B), r2 (A), w2 (A), r2(B), w2(B)
- D. r1(A), w1(A), r2(B), w2(B), r2(A), w2(A), r1 (B), w1 (B)
- E. All (A)-(D) contain conflicts

Question 18 [4 marks] The write-ahead logging (WAL) protocol simply means that:

- A. writing of a data item should be done ahead of any logging operation.
- B. the log record for an operation should be written before the actual data is written.
- C. all log records should be written before a new transaction begins execution.
- D. the log never needs to be written to disk.

Question 19 [4 marks] If a database system supports ACID properties for transaction execution, which of the following pairs of values is a possible result for A and B, after executing the below transactions T1 and T2 concurrently, with an initial value of A=100 and B=200?

T1:	<div>read(B) B=B+50 write(B) read(A) A=A-50 write(A)</div>	T2:	<div>read(B) tmp=B*0.1 B=B-tmp write(B) read(A) A=A+tmp write(A)</div>
------------	--	------------	--

- A. A=70 , B=230
- B. A=50 , B=180
- C. A=120 , B=250
- D. A=50 , B=250
- E. None of the above

Question 20 [4 marks] If a steal/no-force buffer management policy is in place, which of the following is true about system recovery?

- A. Both the Redo and Undo operations are needed
- B. Neither the Redo nor the Undo operations is needed
- C. Redo is needed but Undo is not needed
- D. Redo is not needed but Undo is needed

Question 21 [5 marks]:

You are given the following tables:

`Student(StudId, Name, Addr, Status)`

`Transcript(Id, CrsCode, Semester, Grade)`

Consider a query that outputs the names of all students who took INFS2200 in 2017. An execution plan for that query can be expressed as follows:


$\Pi_{Name}(\sigma_{Id=StudId \text{ AND } CrsCode='INFS2200' \text{ AND } Semester='2017'}(Student \times Transcript))$

In the following, fill in the missing subscripts of the different operators (i.e., Π , σ , \Join) so that to create an optimized plan that is equivalent to the one above.

$\Pi_{NAME} [$
 $(\Pi_{\text{-----}} Student) \Join_{\text{-----}}$
 $(\sigma_{\text{-----}} (\Pi_{\text{-----}} Transcript))$
 $]$

Question 22 [5 marks]::

Consider the following sequences of actions, listed in the order they are submitted by transactions T1, T2, and T3.



Describe how **strict 2PL** with deadlock detection (assume wait-for-graph is used) executes this sequence of actions. Specifically, complete the sequence listed below to show the **lock** and **unlock** requests made by these transactions as well as the blocked and unblocked operations. If a transaction is blocked, assume that all its actions are queued until it is resumed; the DBMS continues with the next action of an unblocked transaction.

T1 acquires shared-lock on A.

T2 blocks waiting for an exclusive-lock on A.

...

Question 23:

For each of the following schedules: 1) construct a precedence graph, 2) determine if the schedule is conflict serializable, and 3) determine the equivalent serial schedule.

(a) **[3 marks]** $r_1(X); r_3(X); w_1(X); r_2(X); w_3(X)$

(b) **[3 marks]** $r_3(X); r_2(X); w_3(X); r_1(X); w_1(X)$

Question 24:

Consider the join $R \bowtie_{R.a=S.b} S$, given the following information about the relations to be joined. The cost metric is the number of block read operations.

Relation R contains 10,000 tuples and has 10 tuples per block.

Relation S contains 2000 tuples and also has 10 tuples per block.

Attribute b of relation S is the **primary** key for S.

Both relations are stored as simple heap files.

(a) [2 marks] What is the cost of joining R and S using a **page-oriented nested loop** join?

(b) [2 marks] How many tuples does the join of R and S produce, at most? and how many blocks are required to store the result of the join back on disk?

END OF EXAMINATION